

Azure Load Balancer

Azure Load Balancer is a Layer-4 network load balancing service that distributes incoming TCP/UDP traffic across multiple backend virtual machines to provide high availability and scalability. It uses frontend IPs, backend pools, health probes, and load-balancing rules to route traffic only to healthy instances.

Main components

→ Frontend IP Configuration

Entry point for traffic.

- Public IP → Internet-facing apps
- Private IP → Internal apps (inside VNet)

→ Backend Pool

A group of resources that receive traffic:

- Virtual Machines
- VM Scale Set instances
- NICs / Private IPs

These are your **actual application servers**.

→ Health Probes

Continuously checks backend servers (TCP/HTTP):

- If VM is **healthy** → traffic continues
- If VM is **unhealthy** → traffic stops to that VM

This gives **automatic failover**.

→ Load Balancing Rules

Define:

- Frontend port (example: 80)
- Backend port (example: 80)
- Protocol (TCP/UDP)
- Which backend pool to use

Types of Load Balancers

→ public Load balancer

- Has public ip
- Used for internet-facing apps
- Ex: Web Servers

→ Private Load balancere

- Uses private IP
- Works only inside Vnet
- Example: app tier → db tier

Steps to create ALB

→ Lets Create Azure Load Balancer

→ Create two Vms with in a vnet with different Subnets

→ Creating a WebVm1

The screenshot shows the 'Create a virtual machine' wizard in the Azure portal. The 'Project details' step is selected. Under 'Subscription', 'Azure subscription 1' is chosen. Under 'Resource group', 'Sai-rg' is selected. In the 'Instance details' section, the 'Virtual machine name' is 'WebVm1', 'Region' is '(Canada) Canada Central', 'Availability options' is 'No infrastructure redundancy required', 'Security type' is 'Standard', and 'Image' is 'Ubuntu Server 24.04 LTS - x64 Gen2 (free services eligible)'. At the bottom, there are buttons for '< Previous', 'Next : Disks >', and 'Review + create'.

→ Adding Vnet and Subnet for WebVm1

The screenshot shows the 'Create a virtual machine' wizard in the Azure portal. The 'Network interface' step is selected. Under 'Virtual network', '(New) vnet-1 (Sai-rg)' is chosen. Under 'Subnet', '(New) snet-1' is selected, with the note '172.16.0.0 - 172.16.0.255 (256 addresses)'. Under 'Public IP', '(new) WebVm1-ip' is selected. Under 'NIC network security group', 'Basic' is selected. Under 'Public inbound ports', 'Allow selected ports' is selected with 'HTTP (80), SSH (22)' chosen. A warning message at the bottom states: '⚠ This will allow all IP addresses to access your virtual machine. This is only recommended for testing. Use the Advanced controls in the Networking tab to'. At the bottom, there are buttons for '< Previous', 'Next : Management >', and 'Review + create'.

→ Creating Second WebVm as WebVm2

The screenshot shows the 'Create a virtual machine' wizard in the Microsoft Azure portal. The 'Basics' tab is selected. In the 'Project details' section, the subscription is set to 'Azure subscription 1' and the resource group is 'Sai-rg'. In the 'Instance details' section, the virtual machine name is 'WebVm2' and it is deployed to the '(Canada) Canada Central' region. A note at the bottom left says: 'This subscription may not be eligible to deploy VMs of certain sizes in certain regions.' At the bottom, there are buttons for '< Previous', 'Next : Disks >', and 'Review + create'.

→ Adding Vnet and Subnet for WebVm2

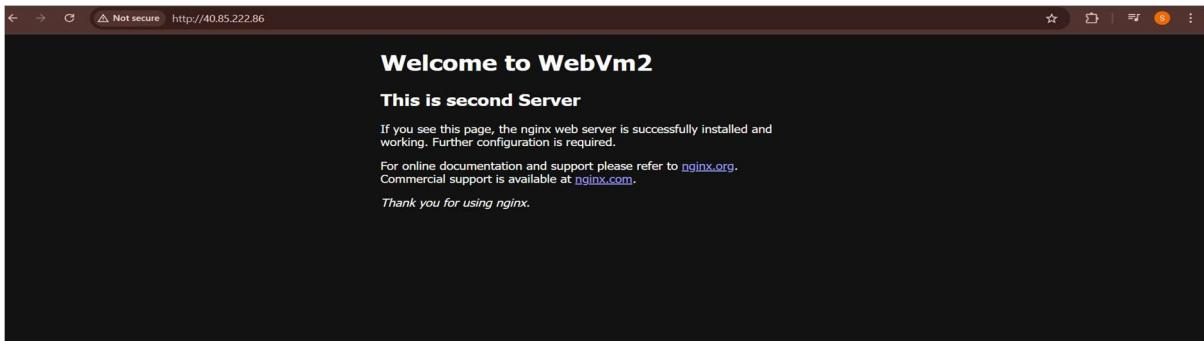
The screenshot shows the 'Network interface' step of the 'Create a virtual machine' wizard. It includes fields for 'Virtual network' (set to 'vnet-1 (Sai-rg)'), 'Subnet' (set to 'snet-2'), 'Public IP' (set to '(new) WebVm2-ip'), 'NIC network security group' (set to 'Basic'), and 'Public inbound ports' (set to 'HTTP (80), SSH (22)'). Below these, a note says: 'Public IP addresses have a nominal charge. [Estimate price](#)'.

→ Then install nginx on two Vms and check it is working or not

→ WebVm1

The screenshot shows a web browser displaying the 'Welcome to WebVm1' page. The page content reads: 'Welcome to WebVm1', 'This is first Server', 'If you see this page, the nginx web server is successfully installed and working. Further configuration is required.', 'For online documentation and support please refer to [nginx.org](#). Commercial support is available at [nginx.com](#).', and 'Thank you for using nginx.'.

→ WebVm2



→ Search for load balancer and click on create and select standard load balancer

→ Creating a Load balancer

Azure load balancer is a layer 4 load balancer that distributes incoming traffic among healthy virtual machine instances. Load balancers uses a hash-based distribution algorithm. By default, it uses a 5-tuple (source IP, source port, destination IP, destination port, protocol type) hash to map traffic to available servers. Load balancers can either be internet-facing where it is accessible via public IP addresses, or internal where it is only accessible from a virtual network. Azure load balancers also support Network Address Translation (NAT) to route traffic between public and private IP addresses. [Learn more](#).

Project details

Subscription * Azure subscription 1

Resource group * Sai-rg

Instance details

Name * PublicLoadBalancer

Region * Central India

SKU * Standard (Distribute traffic to backend resources)

Type * Public

Tier * Regional

[Review + create](#) [Next : Frontend IP configuration >](#) [Download a template for automation](#) [Give feedback](#)

→ Adding Frontend ip to the load balancer

Add frontend IP configuration

PublicLoadBalancer

Name * frontendip

IP version IPv4

IP type IP address

Public IP address * Select public IP address

Add a public IP address

Name * azlbpip

SKU Standard

Tier Regional

Assignment Static

Availability zone * Zone-redundant

Routing preference Microsoft network

[Review + create](#) [Next : Backend pools >](#) [Download a template for automation](#) [Give feedback](#)

→ Adding a backend pool to the load balancer and the two vms for this load balancer

The screenshot shows the 'Add IP configurations to backend pool' dialog in the Azure portal. The 'Resource Name' dropdown is set to 'mybackendpool'. The 'Location' is 'canadacentral' and the 'Virtual network' is 'vnet-1 (Sai-rg)'. Under 'IP configurations', there are two entries: 'WebVm1' and 'WebVm2', both selected. The 'Save' button is at the bottom left, and 'Give feedback' is at the bottom right.

→ Add Load Balancing rule

The screenshot shows the 'Create load balancer' dialog in the Azure portal. The 'Inbound rules' tab is selected. On the right, the 'Add load balancing rule' dialog is open, showing the following configuration:

- IP version:** IPv4 (selected)
- Frontend IP address:** frontendip (To be created)
- Backend pool:** mybackendpool
- Protocol:** TCP (selected)
- Port:** 80
- Backend port:** 80
- Health probe:** No existing probes (Create new)
- Session persistence:** None

Buttons at the bottom include 'Review + create', '< Previous', 'Next : Outbound rule >', 'Download a template for automation', and 'Give feedback'.

→ Adding Health probe

The screenshot shows the 'Create health probe' dialog in the Azure portal. A tooltip explains: 'Health probes are used to check the status of a backend pool instance. If the health probe fails to get a response from a backend instance then no new connections will be sent to that backend instance until the health probe succeeds again.' The dialog fields are:

- Name:** my probe
- Protocol:** TCP
- Port:** 80
- Interval (seconds):** 5
- Used by:** Not used

Buttons at the bottom are 'Save' and 'Cancel'.

→ Check the review and click on create

→ The frontend ip is the public ip of your load balancer

→ Copy and search in browser

The screenshot shows the Microsoft Azure Load Balancer Overview page. It displays various configuration details such as Resource group (Sai-rg), Location (Canada Central), Subscription (Azure subscription 1), SKU (Standard), Tier (Regional), and Backend pool (mybackendpool). The Frontend IP address is listed as 20.220.152.87 (azbip). A 'Copy to clipboard' button is visible next to the IP address.

→ First I search in this browser the traffic route to my webVm2

The screenshot shows a browser window displaying the 'Welcome to WebVm2' page. The page content includes 'Welcome to WebVm2', 'This is second Server', and a note: 'If you see this page, the nginx web server is successfully installed and working. Further configuration is required.' It also mentions 'For online documentation and support please refer to [nginx.org](#). Commercial support is available at [nginx.com](#)'. The footer says 'Thank you for using nginx.'

→ When I stop my WebVm2

The screenshot shows the Microsoft Azure Compute Infrastructure Virtual machines page. It lists several virtual machines, including vm01, vm02, WebVm1, and WebVm2. A confirmation dialog box is overlaid, asking 'Do you want to stop all the selected virtual machines?' with 'Yes' and 'No' buttons. Below the list, there is a table showing columns for Name, Subscription, Resource Group, Location, Status, Operating system, Size, Public IP address, and Disks.

→ Then traffic routes to my WebVm1

The screenshot shows a browser window displaying the 'Welcome to WebVm1' page. The page content includes 'Welcome to WebVm1', 'This is first Server', and a note: 'If you see this page, the nginx web server is successfully installed and working. Further configuration is required.' It also mentions 'For online documentation and support please refer to [nginx.org](#). Commercial support is available at [nginx.com](#)'. The footer says 'Thank you for using nginx.'

→ When you start WebVm2 and search in another browser you can see the webvm2 after some time.